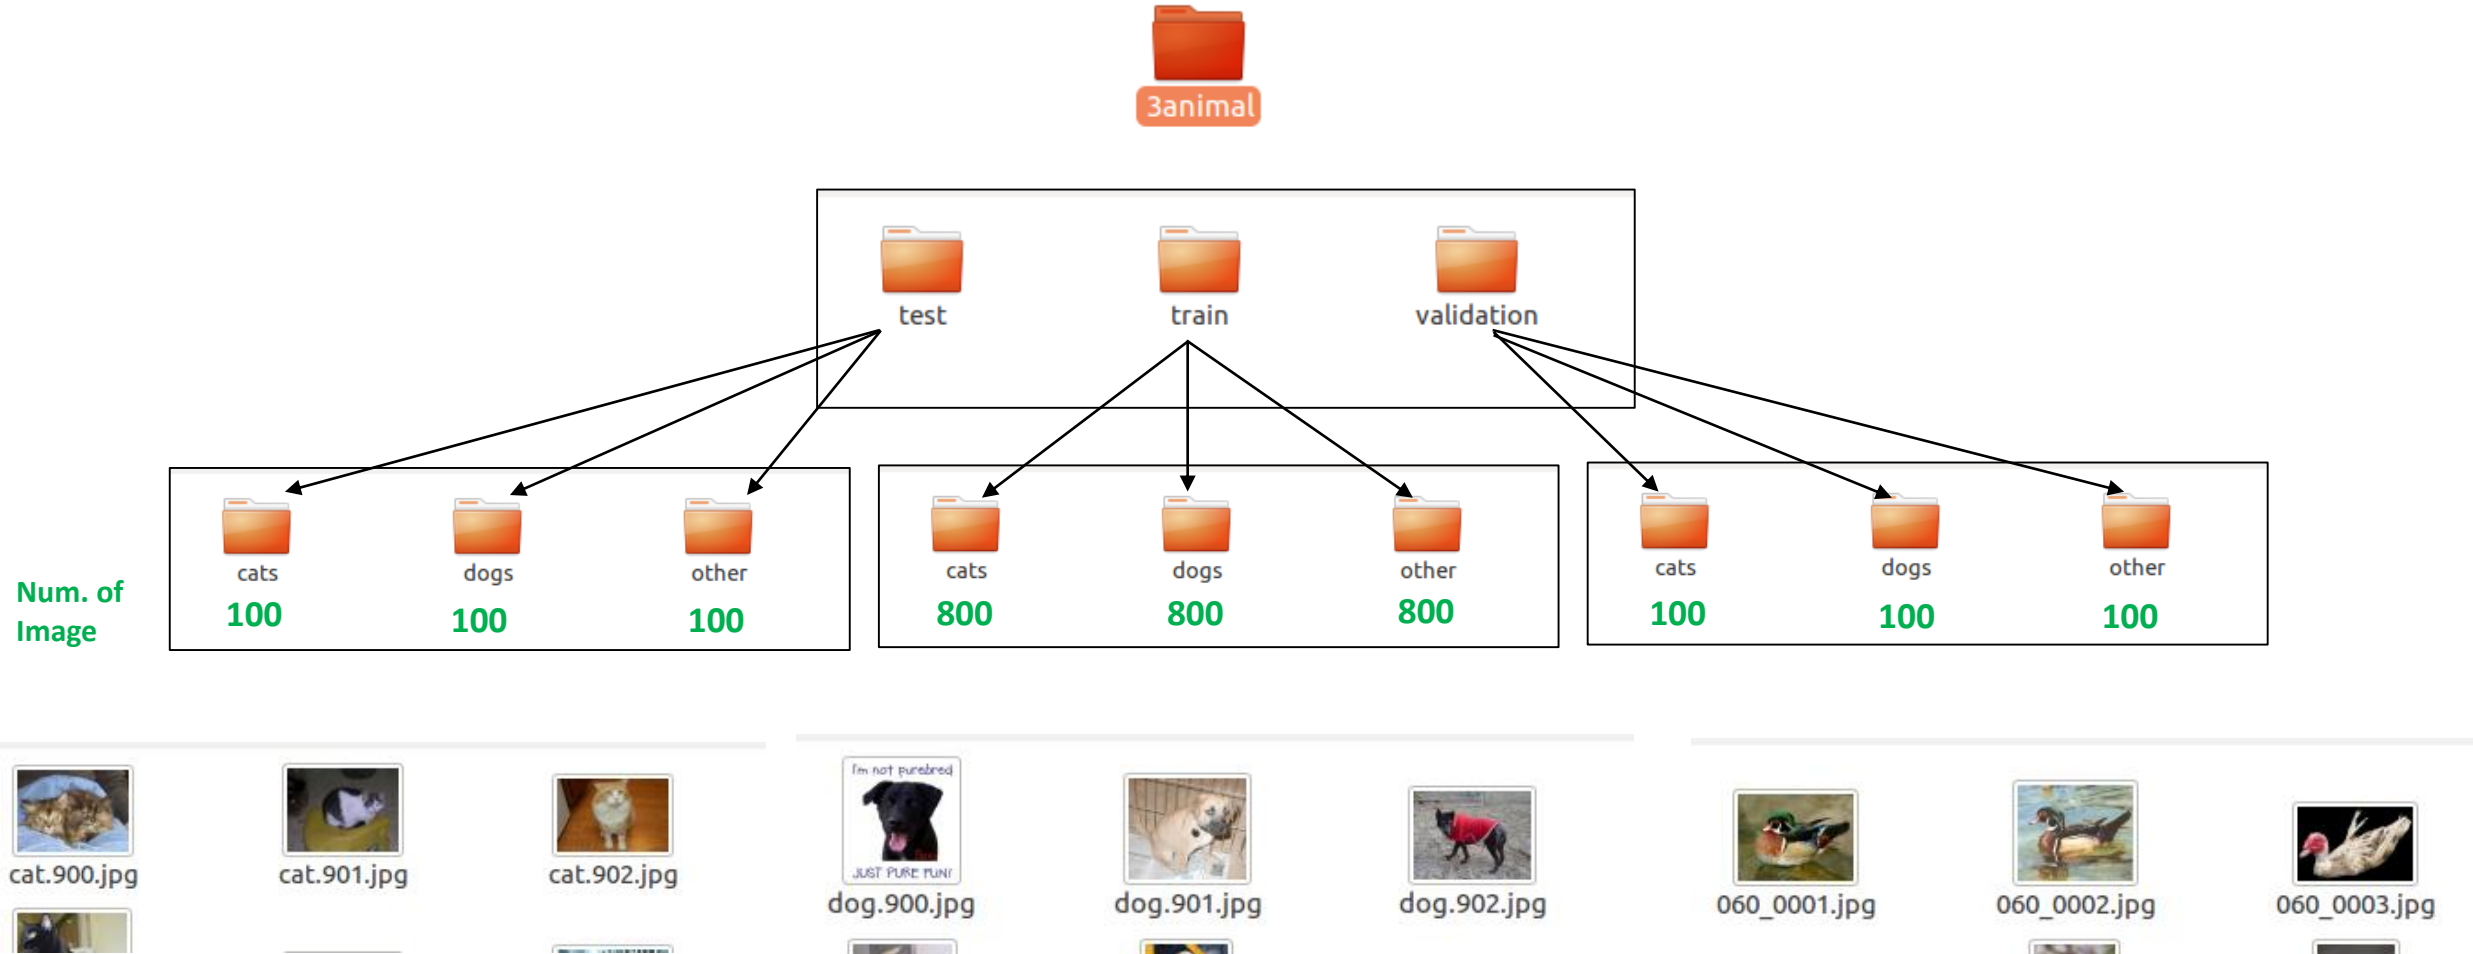


# DOE for Weed Detection Simulation

Sayem

3/1/2018


# DATA



## Steps

1. Build a small CNN (4 layers)
2. Small CNN + data Augmentation
3. Use a Pre-trained CNN
4. Use Pre-trained CNN + data augmentation
5. Use Pre-trained CNN + data augmentation + train last few layers of Pre-trained CNN
6. Hyperparameter tuning

# 1. Build a small CNN (4 layers)





Layer (type)	Output Shape
conv2d_1 (Conv2D)	(None, 148, 148, 32)
max_pooling2d_1 (MaxPooling2)	(None, 74, 74, 32)
conv2d_2 (Conv2D)	(None, 72, 72, 64)
max_pooling2d_2 (MaxPooling2)	(None, 36, 36, 64)
conv2d_3 (Conv2D)	(None, 34, 34, 128)
max_pooling2d_3 (MaxPooling2)	(None, 17, 17, 128)
conv2d_4 (Conv2D)	(None, 15, 15, 128)
max_pooling2d_4 (MaxPooling2)	(None, 7, 7, 128)
flatten_1 (Flatten)	(None, 6272)
dense_1 (Dense)	(None, 512)
dense_2 (Dense)	(None, 1)

# 1. Build a small CNN (4 layers)

Run Time (GPU): 46 secs


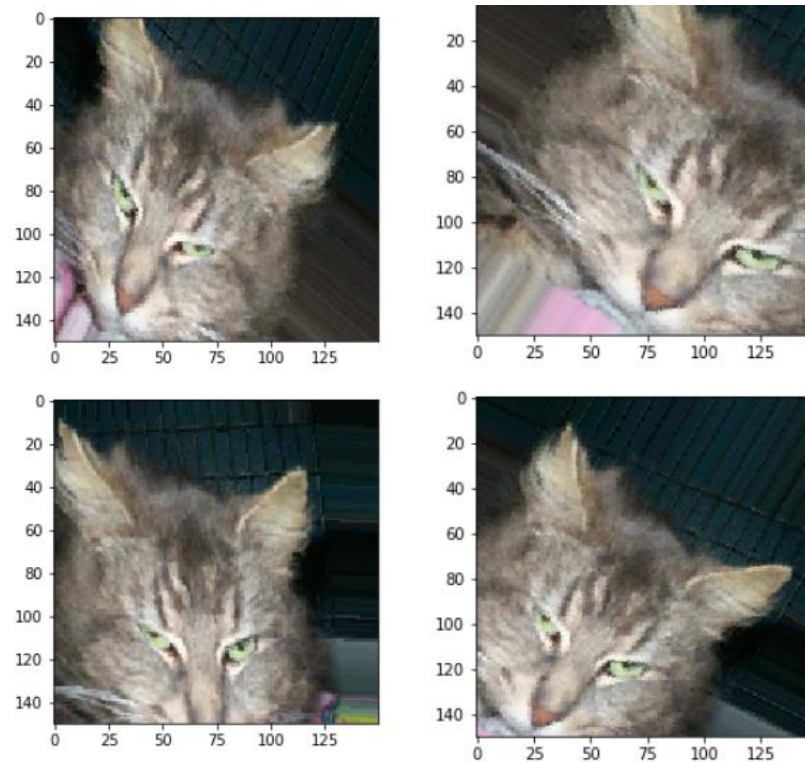
Overfitting : Increase training data to compensate



## 2. Small CNN + data Augmentation

**Data augmentation** takes the approach of generating more training data from existing training samples, by "augmenting" the samples via a number of random transformations that yield believable-looking images.


```
datagen = ImageDataGenerator(  
    rotation_range=40,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest')
```




# ImageNet

- The **ImageNet** project is a large visual [database](#) designed for use in [visual object recognition software](#) research
- Over 14 million URLs of images have been hand-annotated by ImageNet to indicate what objects are pictured
- ImageNet contains over 20 thousand ambiguous categories
- Now, in the case of **top-1** score, you check if the top class (the one having the highest probability) is the same as the target label.
- In the case of **top-5** score, you check if the target label is one of your top 5 predictions (the 5 ones with the highest probabilities).

# 3. Use a Pre-trained CNN



Schematic Diagram of VGG16 Model:



Schematic Diagram of the 27-layer Inception-V1 Model (Idea similar to that of V3):





# 3. Use a Pre-trained CNN

1. **Chollet, François.** "Xception: Deep learning with depthwise separable convolutions." *arXiv preprint (2016)*.
2. **Simonyan, Karen, and Andrew Zisserman.** "Very deep convolutional networks for large-scale image recognition." *arXiv preprint arXiv:1409.1556 (2014)*.
3. **Szegedy, Christian, et al.** "Inception-v4, inception-resnet and the impact of residual connections on learning." *AAAI. Vol. 4. 2017*.
4. **Huang, Gao, et al.** "Densely connected convolutional networks." *Proceedings of the IEEE conference on computer vision and pattern recognition. Vol. 1. No. 2. 2017*.

## Available Pre-trained CNN models

Model	Size	Top-1 Accuracy	Top-5 Accuracy	Parameters	Depth
Xception	88 MB	0.790	0.945	22,910,480	126
VGG16	528 MB	0.715	0.901	138,357,544	23
VGG19	549 MB	0.727	0.910	143,667,240	26
ResNet50	99 MB	0.759	0.929	25,636,712	168
InceptionV3	92 MB	0.788	0.944	23,851,784	159
InceptionResNetV2	215 MB	0.804	0.953	55,873,736	572
MobileNet	17 MB	0.665	0.871	4,253,864	88
DenseNet121	33 MB	0.745	0.918	8,062,504	121
DenseNet169	57 MB	0.759	0.928	14,307,880	169
DenseNet201	80 MB	0.770	0.933	20,242,984	201

**Xception** input size for this model is 299x299 or lower


**VGG16** input size for this model is 224x224 or lower

**InceptionResNetV2** input size for this model is 299x299 or lower



**DenseNet201** input shape has to be (224, 224, 3)

# 3. Use a Pre-trained CNN


**VGG16**  
i5 cpu: 13 min 58 sec  
GPU : 44 sec



**InceptionResNetV2**  
i5 cpu: 9 min 44 sec  
GPU : 1 min 22 sec




**Xception**  
i5 cpu: 13 min 44 sec  
GPU : 1 min 40 sec




**DenseNet201**  
i5 cpu: 8 min 59 sec  
GPU : 1 min 24 sec

# 4. Use Pre-trained CNN + data augmentation


**VGG16**  
**GPU : 8 min 46 sec**




**Xception**  
**GPU : 8 min 40 sec**



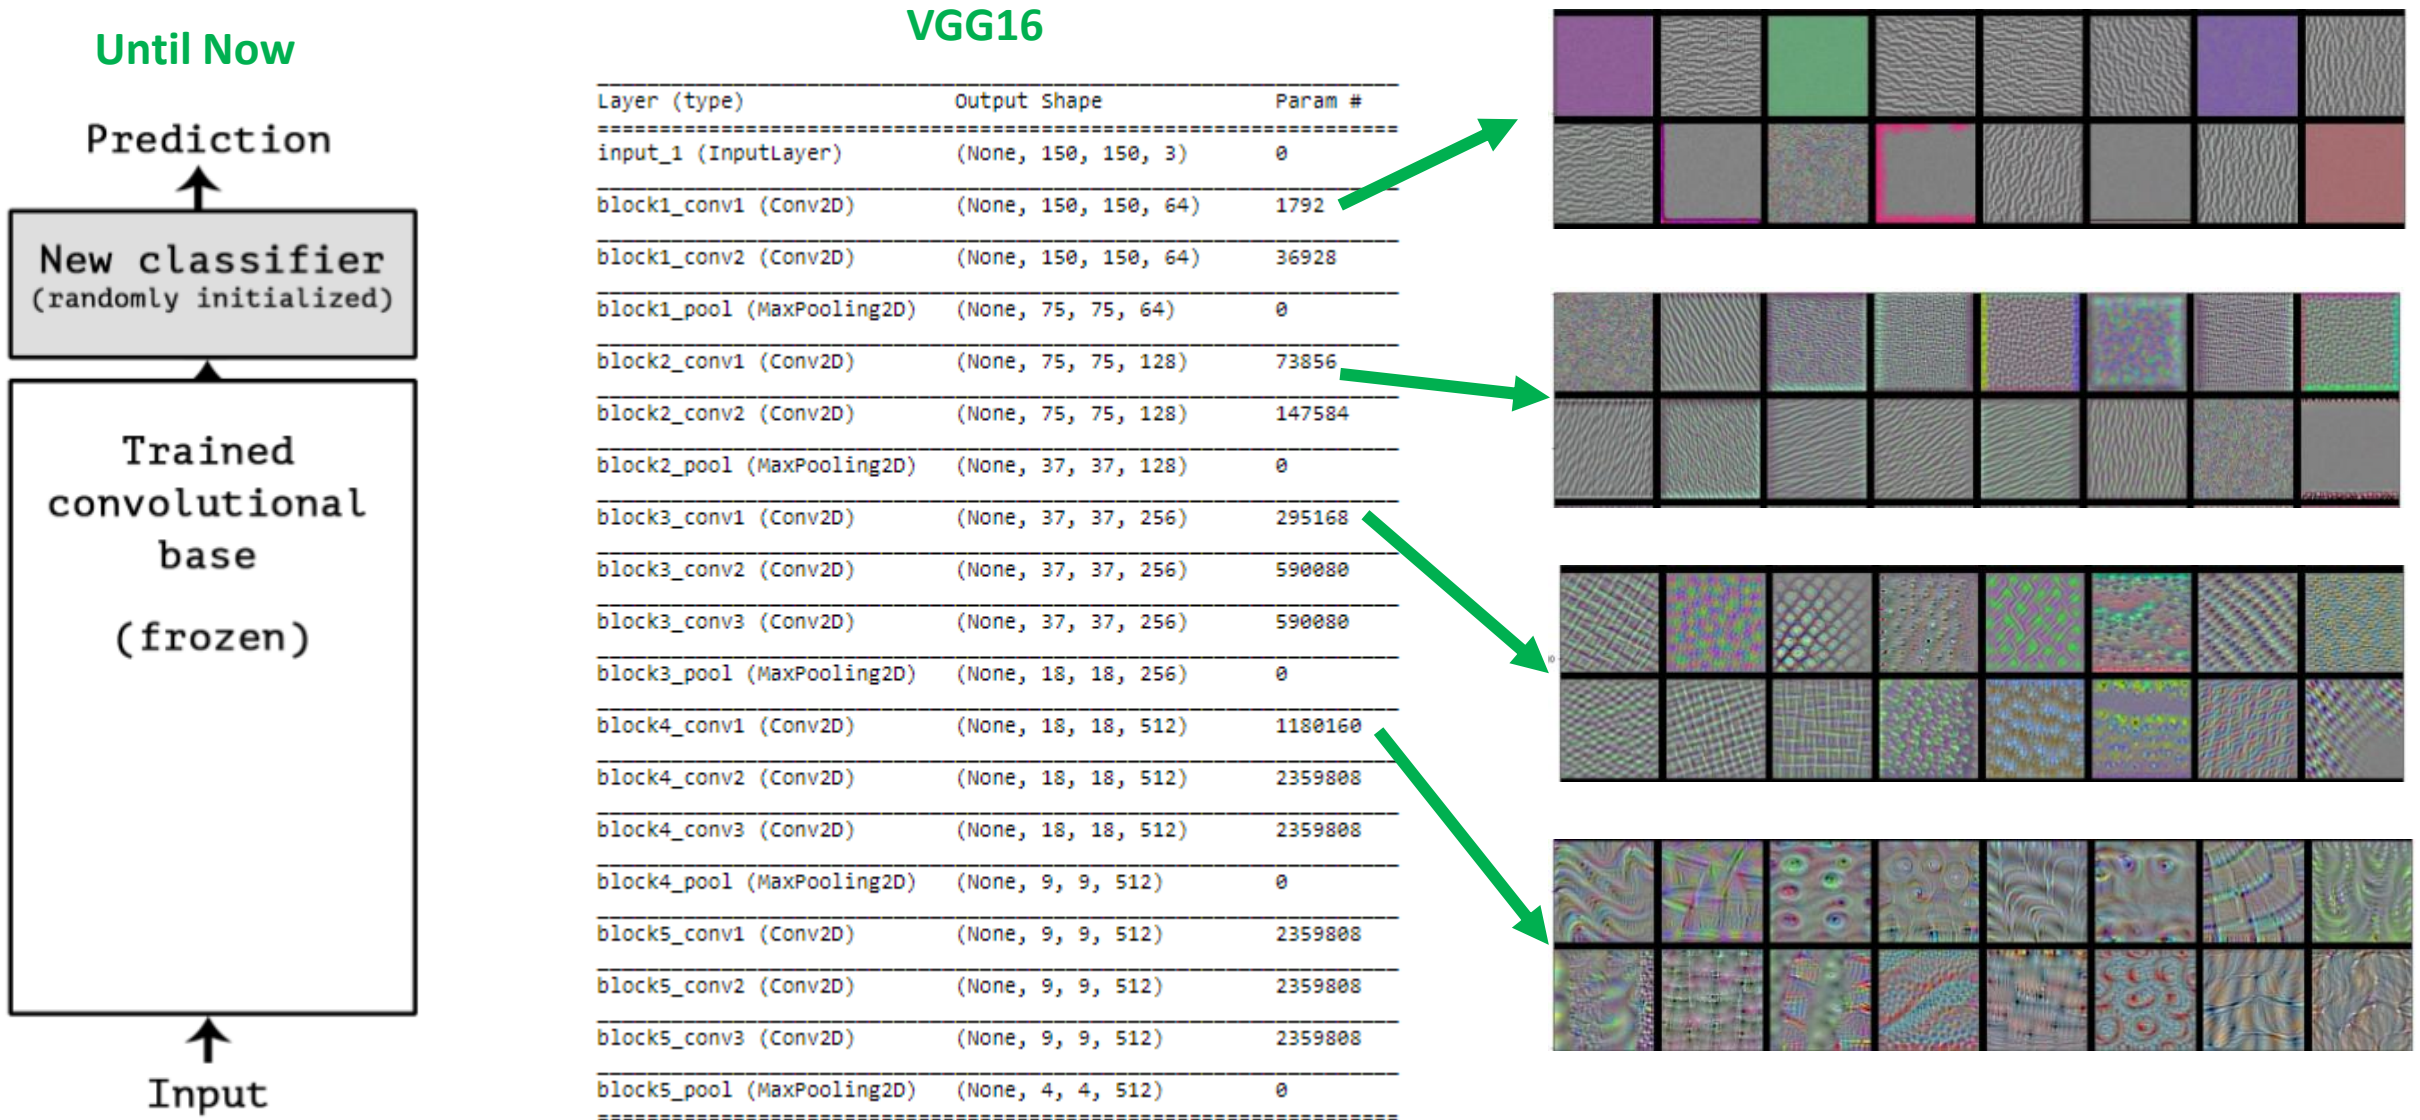
**InceptionResNetV2**  
**GPU : 9 min 55 sec**



**DenseNet201**  
**GPU : 19 min 56 sec**



# 5. Use Pre-trained CNN + data augmentation + train last few layers of Pre-trained CNN





# 5. Use Pre-trained CNN + data augmentation + train last few layers of Pre-trained CNN




These filter visualizations tell us a lot about how convnet layers see the world: each layer in a convnet simply learns a collection of filters such that their inputs can be expressed as a combination of the filters. The filters in these convnet filter banks get increasingly complex and refined as we go higher-up in the model:

- The filters from the first layer in the model (block1\_conv1) encode simple directional edges and colors (or colored edges in some cases).
- The filters from block2\_conv1 encode simple textures made from combinations of edges and colors.
- The filters in higher-up layers start resembling textures found in natural images: feathers, eyes, leaves, etc


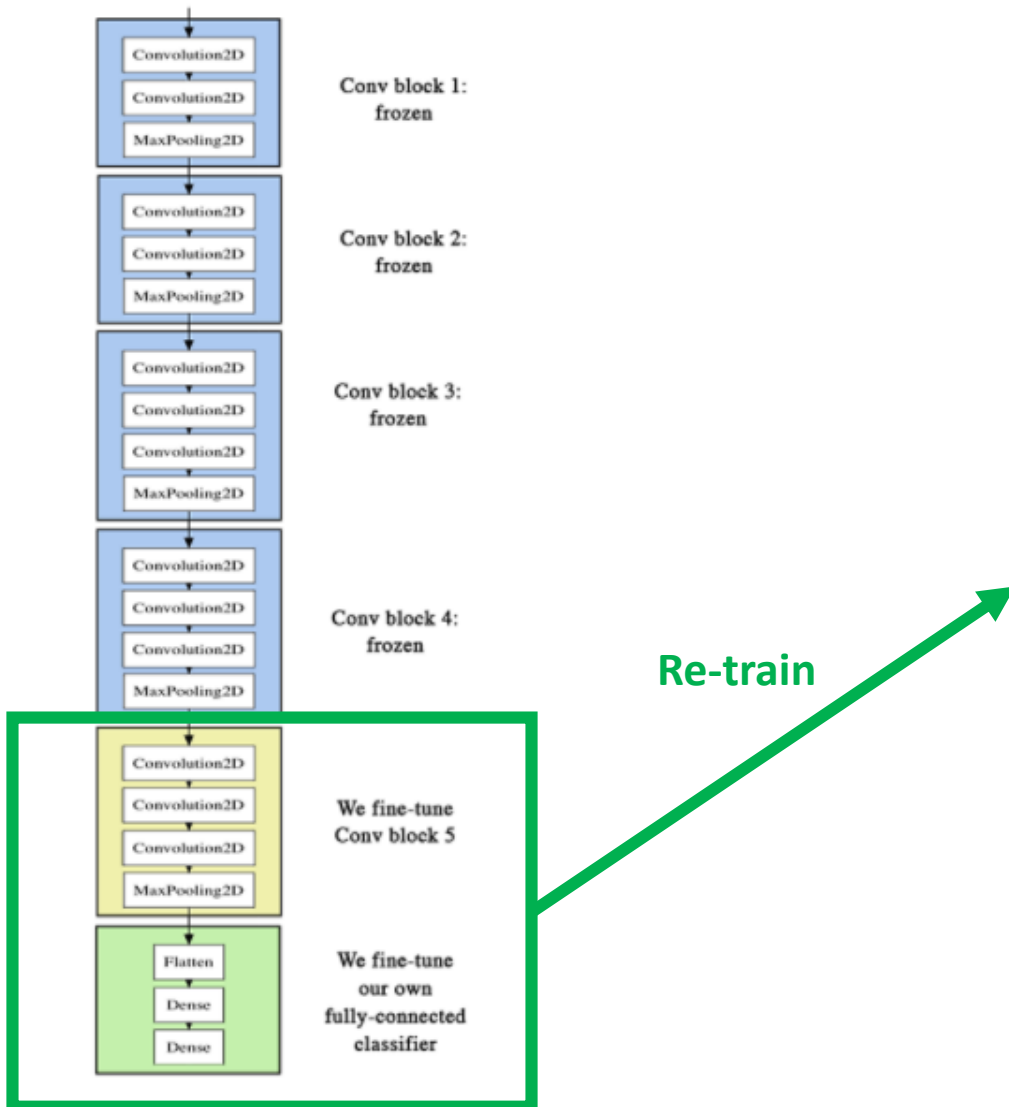
## VGG16

```
Layer (type)
=====
input_1 (InputLayer)
block1_conv1 (Conv2D)
block1_conv2 (Conv2D)
block1_pool (MaxPooling2D)
block2_conv1 (Conv2D)
block2_conv2 (Conv2D)
block2_pool (MaxPooling2D)
block3_conv1 (Conv2D)
block3_conv2 (Conv2D)
block3_conv3 (Conv2D)
block3_pool (MaxPooling2D)
block4_conv1 (Conv2D)
block4_conv2 (Conv2D)
block4_conv3 (Conv2D)
block4_pool (MaxPooling2D)
block5_conv1 (Conv2D)
block5_conv2 (Conv2D)
block5_conv3 (Conv2D)
block5_pool (MaxPooling2D)
```

Filter



# 5. Use Pre-trained CNN + data augmentation + train last few layers of Pre-trained CNN



Run time; 29 min 50 sec

## 6. Hyperparameter tuning

- How many **layers** should you stack? How many **units** or **filters** should go in each layer? Should you use 'relu' as **activation**, or a different function? Should you use **BatchNormalization** after a given layer? How much **dropout** should you use?....etc
- Mostly applicable for NN which are built from scratch.
- Have to look into tools available.

# Things to do:

- Use our custom “trained classifier” to detect object(weed) from video feed.
- Check maximum FPS.
- Create weed dataset + Interlink video feed with ROS (Khunsa & Harsha).